



# State-of-the-Art Image Tampering Detection in Car Damage Claims (Europe & Germany)

## Market Overview: Fraud Prevalence & Impact (Germany vs Europe)

**Prevalence & Cost of Fraud:** Insurance fraud is a significant issue across Europe. Estimates indicate that about **10% of all insurance claim costs in Europe are tied to fraud** (detected or undetected) <sup>1</sup>. In Germany, fraud levels are similar – industry experts estimate roughly **one in ten claims is dubious** <sup>2</sup>. The financial impact is substantial: German insurers in 2024 expected over **€6 billion in losses from fraudulent claims** (up from ~€4–5 billion in prior years) <sup>3</sup>. Notably, **motor claims make up about half** of these dubious cases <sup>4</sup>, reflecting how common car-related fraud is. This fraud burden ultimately drives up premiums and costs for honest customers, as seen in the UK where rising fraud (including fake crash photos) contributed to a 33% annual jump in motor insurance prices <sup>5</sup>.

**Rising Digital Manipulation Threat:** A key driver for the growing fraud problem is the **ease of digital image manipulation**. Both simple “**shallowfakes**” (edited with common photo apps) and AI-generated “**deepfakes**” are on the rise, enabling fraudsters to fabricate convincing accident evidence entirely on a computer <sup>6</sup> <sup>7</sup>. German insurers warn that **new AI tools lower the barrier** for fraud – making it trivial to alter or generate damage photos <sup>8</sup>. Allianz (UK) reported a **300% increase in incidents of digitally doctored images** in just one year (2021–22 to 2022–23) <sup>9</sup>. Zurich UK likewise sees manipulated crash photos as “one of the most emerging threats” in motor fraud <sup>6</sup>. In short, the **credibility of claim photos can no longer be taken at face value**, pressuring insurers to respond with equally sophisticated detection measures.

**Market Growth & Adoption:** In response, the market for fraud detection technology – particularly AI-driven solutions – is booming. Globally, the **insurance fraud detection solutions market** (spanning software and services) reached an estimated \$7.5 billion in 2024 and is projected to grow ~20–25% annually, exceeding \$9 billion in 2025 <sup>10</sup> and over \$22 billion by 2029 <sup>11</sup>. Europe is a major part of this trend: many European insurers are investing in AI for claims fraud detection (e.g. a recent survey found about **50% of non-life insurers in Europe already using AI in various processes** including fraud screening) <sup>12</sup>. **Germany's market** is seeing new specialized startups and pilot programs. For example, Berlin-based insurtech **VAAHRAFT launched in 2023** and by mid-2024 had its **first 3 paid pilots with German insurers (and 7 more in pipeline)** <sup>13</sup>, reflecting strong demand for image-fraud solutions in the German market. Analysts call the outlook “promising,” with rapid adoption of AI/analytics expected among both large and mid-sized German insurers <sup>14</sup> <sup>15</sup>. In summary, **European insurers are actively scaling up anti-fraud tech** to keep pace with tech-savvy fraudsters, creating a fast-growing market for image tampering detection tools.

## Common Image Fraud Schemes in Car Damage Claims

*Example of a tampered damage photo:* The left image is a genuine van photo lifted from social media, and the right image is a **fake “damaged” version** with cracks digitally added to the bumper <sup>16</sup>. Fraudsters use such techniques to fabricate accident evidence.

Modern car insurance fraud involving images typically falls into a few categories:

- **Photoshopped Damage:** Using editing tools (Photoshop or even mobile apps) to **add, remove, or exaggerate damage** on a vehicle. For instance, scammers might clone existing dents/scratches to multiple spots, add fake cracks or smash effects, or hide telltale signs that the photo is old <sup>17</sup> <sup>18</sup>. The above van example illustrates how easily a cracked bumper can be forged. These are sometimes called "*shallowfakes*" – simple edits that don't require AI <sup>19</sup>.
- **Image Reuse & Misattribution:** Reusing an **old or unrelated photo** to support a new claim. A fraudster might submit damage photos from a previous incident or steal images from the internet (e.g. from salvage auction sites or social media) and claim them as their own loss <sup>18</sup> <sup>20</sup>. In one case, scammers took a van photo off Facebook and filed a bogus claim for a fake crash using that image <sup>21</sup>. Others have taken images of totaled vehicles online and **edited the license plate** to match their insured car, then claimed a total loss incident that never occurred <sup>20</sup>. Without tools, a claims handler might not realize the photo was recycled from a different context.
- **AI-Generated Images:** Using **generative AI** (deepfakes) to create wholly fake but photorealistic accident photos. For example, a fraudster could generate an image of a smashed car or hail damage that never actually happened. These AI-created images are **highly realistic and coherent**, containing no obvious splicing seams. This makes them harder to detect with traditional forensic tricks. While still emerging, such deepfakes are a growing concern – experts warn that fraudsters can now concoct an entire fake accident from behind a keyboard, with computer-generated crash photos, fake repair bills, and even AI-generated "proof" documents <sup>7</sup>.
- **Context Manipulation:** Simply **cropping or altering image context** to mislead. For instance, cropping out a date stamp, or using angles/lighting to conceal that damage existed pre-incident. Some will adjust timestamps, remove metadata, or screenshot images (to strip metadata) before submission <sup>22</sup>. Others might overlay elements – e.g. merging two photos – to fabricate a scene (such as superimposing an object into a photo to claim it was damaged/stolen) <sup>23</sup>. These subtler tweaks can fool manual reviewers, especially if they're rushed.

**Impact on Insurers:** These tactics lead to **inflated payouts and bogus claims**, directly hitting insurers' loss ratios. They also **erode trust** – legitimate claimants might face more scrutiny or delays because insurers must double-check image evidence. Operationally, the rise of tampered images forces insurers to spend more on investigations and slows down straight-through processing. In Germany and Europe where digital claims (photo-based e-claims) are encouraged for efficiency, this is a double-edged sword: it speeds up honest claims but also opens a door for "fraud at scale" (one person can digitally mass-produce fake claims). This has made **automated image verification essential** to maintain both efficiency and integrity in claims handling <sup>24</sup> <sup>25</sup>.

## Detection Techniques: How Technology (Especially AI/ML) Fights Image Fraud

Modern image tampering detection for insurance relies on a **multi-layered, AI-driven approach**. Key techniques include:

- **Metadata & EXIF Analysis:** The first line of defense is checking the **image's metadata** for inconsistencies <sup>26</sup>. This includes verifying if the **timestamp** aligns with the reported accident

time, if the **GPS location** matches the claimed location of loss, and whether the **device/camera model** is plausible (e.g. matches the policyholder's phone or at least isn't an obvious stock photo signature) <sup>27</sup>. Metadata can also reveal editing – for example, some files show software tags or indicate a copy/download. If an image's embedded data shows it was created long before the accident, or edited with a graphics program, that's a red flag <sup>18</sup>. *Trade-off:* metadata checks are fast and cheap, but **fraudsters can strip or alter metadata easily** (e.g. by re-saving or screenshotting the image) <sup>22</sup>. So, metadata anomalies are suggestive but not conclusive – many genuine claim photos might lack metadata due to innocuous reasons (privacy settings, messaging apps stripping data, etc.), so context is important to avoid false alarms.

- **Visual Forensics (Pixel-Level Analysis):** Advanced tools perform **pixel-level scrutiny** of images to spot signs of manipulation invisible to the naked eye. One classic method is **Error Level Analysis (ELA)**, which highlights compression artifacts – if parts of an image were edited or spliced from another image, they often show different error levels after recompression <sup>28</sup>. Other forensic analyses look at lighting inconsistencies (e.g. shadows at wrong angle), noise patterns, or anomalies in reflections and edges <sup>29</sup> <sup>30</sup>. For example, duplicated scratch marks or slight blurring around an inserted object can indicate tampering. Some solutions analyze Photo Response Non-Uniformity (PRNU), essentially the “fingerprint” noise of a camera, to see if the photo is an original from a device or a digitally altered composite. **Visual anomalies** like perfectly identical damage patterns, mismatched shadows, or abrupt changes in pixel noise can be flagged automatically by software <sup>29</sup> <sup>30</sup>. However, as **fully AI-generated images have become more coherent**, methods like ELA are less effective since a single AI-generated photo may not contain spliced parts – it's internally consistent <sup>31</sup>. This is where AI/ML steps up the game.
- **AI/ML Image Classification:** Machine learning, especially deep learning, is now central to tampering detection. Insurers and tech providers use **computer vision models trained to distinguish real vs fake images**. For instance, **Vision Transformers (ViT)** and convolutional neural networks can be trained on large datasets of known fake images (including AI-generated fakes and Photoshopped examples) and genuine claim photos <sup>32</sup>. These models learn subtle statistical differences – e.g. slight spectral artifacts or texture patterns – that human eyes can't see. According to Shift Technology, **transformer-based models have achieved impressive accuracy in detecting AI-generated images** in their tests <sup>32</sup>. AI can also detect things like “*unnatural shadowing or duplicated objects*” that suggest cloning in an image <sup>33</sup> <sup>34</sup>. The big advantage of ML is that it can **adapt and improve**: by learning from new fraud cases, it keeps up with evolving tricks <sup>35</sup>. In fact, some providers claim they can update their deep learning models within 48 hours to respond to novel generative AI outputs <sup>36</sup>. The role of AI here is truly a game of “**fight fire with fire**” – using AI to combat AI-assisted fraud <sup>37</sup> <sup>38</sup>. Without ML, insurers would be hopelessly outpaced by the speed and realism with which fraudsters can now produce fake images.
- **Cross-Image Comparison & Search:** Another technique is to check if the submitted photo has appeared elsewhere. Tools perform **reverse image searches** or compare against databases of past claims. This can catch fraud like the reuse of the same damage photo in multiple claims. For example, a system might flag that a car damage photo was used in a claim last year (possibly with a different insurer). Insurtech solutions in Germany even offer “**cross-insurance duplicate checks**,” meaning they can detect if an image was used in a claim with another company <sup>39</sup>. This kind of collaboration (often via a shared platform or vendor) is invaluable in Europe's fragmented market, since a fraudster might shop the same fake claim to several insurers. Additionally, AI-based image similarity algorithms can identify when two photos are of the *same*

*vehicle damage* even if from different angles or backgrounds, helping expose staged or duplicate claims <sup>40</sup> <sup>41</sup>.

- **Contextual and Semantic Analysis:** Cutting-edge solutions go beyond the image pixels by evaluating the **image in context of the claim**. This involves checking if the **content "makes sense" given the incident details**. For example, if a claim description says a car's left side was scraped in a parking lot, but the photo shows heavy front-end collision damage, a semantic analysis would flag a mismatch. AI vision models can classify the type of damage and even the force direction, then cross-check with the accident report. Some insurers integrate image analysis with policy data – e.g. is the car model and color in the photo the same as on record? Does the weather in the photo match the date/time of loss? These context checks help reduce false positives and pinpoint truly suspicious cases <sup>42</sup> <sup>43</sup>. By combining **multiple signals (metadata, pixel anomalies, claim context)**, modern systems achieve far greater accuracy than any single method alone <sup>44</sup> <sup>45</sup>.

**AI/ML at the Core:** Overall, **AI and machine learning drive most of these advanced detection techniques**. Rules and forensic analyses still play a role, but ML excels at spotting complex patterns and can continually learn from new fraud examples. Crucially, AI can do in **seconds** what a human might miss in hours – one service checks an image's authenticity “in seconds” via API <sup>46</sup>, allowing suspicious claims to be flagged almost in real-time. This speed is important because many insurers aim to process simple claims in minutes; an automated image check has to keep up and not become a bottleneck.

*Trade-offs:* Despite the power of AI, challenges remain. **False positives** can occur – e.g. an AI might flag a perfectly legitimate photo that has low quality or was oddly cropped as “suspect.” Likewise, **false negatives** are possible if a new tampering method isn't yet learned by the model <sup>47</sup>. Maintaining a low false-positive rate is critical; denying or delaying a valid claim due to a detection error can harm customer trust <sup>48</sup>. Thus, best practice is a **hybrid workflow**: AI does the heavy lifting to **screen and score images**, but questionable cases get escalated to human investigators for final judgment <sup>49</sup> <sup>50</sup>. AI essentially triages claims – most go through if clean, while the riskiest get human attention. Continuous training (using feedback from those investigations) then improves the AI. In summary, ML/AI is indispensable in this space for its **speed, scalability, and ability to catch subtle or AI-generated fraud that humans or simple rules would miss**, but it works best in concert with human expertise and broader fraud analytics.

## Solutions Landscape: Commercial Tools vs. In-House Approaches

### Commercial Off-the-Shelf (COTS) Solutions

A number of **COTS solutions** have emerged to help insurers detect image fraud without reinventing the wheel. These range from specialist startups to established insurance tech firms:

- **FRISS “Media Check”:** FRISS (a Dutch-founded insurance fraud platform used globally) offers a **Media Check module** that analyzes claim photos and documents for signs of tampering. It uses a combination of **metadata analysis, anomaly detection, and reuse checks** <sup>51</sup>. In 2025, FRISS partnered with Verisk to incorporate Verisk's advanced **Digital Media Forensics** into this tool, specifically for the European market <sup>52</sup>. The system **scans each submitted photo for anomalies** with AI-driven algorithms, flags manipulated images, and can even **trace file origins** or identify if an image's history is suspicious <sup>53</sup>. FRISS Media Check integrates into the claims workflow (including at First Notice of Loss) to automatically screen out fraudulent evidence **before claims are paid** <sup>54</sup> <sup>45</sup>. It's typically part of a larger fraud detection suite that also looks

at claim patterns and networks. **Trade-off:** As a comprehensive solution, it likely comes with enterprise pricing, but it aims to **save costs by preventing payouts and reducing manual review load**. FRISS emphasizes that tools like this **free up SIU investigators** to focus on truly complex cases while allowing straight-through processing for honest claims <sup>55</sup>.

- **VAARHAFT Fraud Scanner:** VAARHAFT is a German InsurTech startup (based in Berlin) focusing specifically on **image-based insurance fraud**. Launched in 2023, their solution provides an **API and web interface** to verify the credibility of claim photos in real-time <sup>56</sup> <sup>57</sup>. It uses **deep learning and computer vision** to detect anything from copy-paste edits to AI-synthesized images. One notable feature tailored to Germany is a “**cross-insurance duplicate check**” – since VAARHAFT works with multiple insurers, their network can spot if the same image has been used in claims across different companies <sup>39</sup>. The tool also can **visually highlight the areas of an image that appear processed or doctored** for an investigator to review <sup>39</sup>. VAARHAFT touts rapid model updates (within 48 hours for new AI manipulation techniques) and full **GDPR compliance** (no customer data is stored or used to train AI without permission) <sup>46</sup> [28timage]. Being a newer solution, insurers often pilot it alongside existing processes. **Cost model:** As a startup, VAARHAFT likely offers flexible pricing – possibly per image scan or annual license – to encourage adoption. It is marketed as **highly configurable** for each insurer’s needs <sup>46</sup>. The benefit is a cutting-edge, focused tool; the risk might be relying on a young company whose algorithms are still evolving. However, early recognition (innovation awards in Europe <sup>58</sup>) and pilot uptake suggest it’s on the right track for the German market.
- **Shift Technology (Force™):** Shift Technology is a Paris-based AI provider widely used for claim fraud detection. While Shift’s platform looks at the entire claim (data, networks, etc.), it also incorporates **image analysis capabilities**. Shift’s approach uses **Generative AI and predictive analytics** to detect “generated, manipulated, or re-used images” in claims <sup>59</sup>. They leverage **vision AI models and even large language model techniques** (for analyzing documents) in one package. For example, Shift demonstrated identifying a fraud case where the **same refrigerator photo** was used in multiple claims by different people – their AI’s image similarity scoring exposed the reuse <sup>40</sup>. In automotive claims, Shift’s solutions can flag digitally added damage or images that don’t match the claim facts. Their 2025 whitepaper notes that **transformer-based vision models** are delivering the best results for detecting AI-fakes <sup>32</sup>. **Integration:** Shift’s system typically integrates via API or as a SaaS platform that ingests claim info and media, returning fraud alerts. Many large European insurers (and global ones like Tokio Marine) have partnered with Shift for AI-driven fraud screening. **Pricing** is not public, but as an enterprise SaaS, it likely involves an annual subscription or volume-based model (often tied to number of claims processed). The value proposition is **accuracy and breadth** – combining image forensics with all other fraud signals in one AI platform.
- **Attestiv:** Attestiv is a vendor offering **digital media authenticity** solutions, known for tamper-proofing and AI analysis of photos. While based in the US, it targets insurance use-cases like claims. Attestiv’s platform can **detect deepfakes, fake images, and document forgeries**, and also offers a **secure capture app** to create verifiable photos. They highlight ROI strongly – providing tools like an ROI calculator to show insurers how much they could save by catching media fraud up front <sup>60</sup> <sup>61</sup>. Attestiv emphasizes turning fraud prevention into a “savings engine” by **preventing bogus claims from ever entering the system** <sup>62</sup>. Some insurers might use Attestiv’s tech in their mobile apps to ensure photos are original (e.g., using blockchain or hash to detect later tampering) – which is a *preventative* approach. **Cost model:** Likely a SaaS fee based on usage; they position it as easily justifiable by the fraud losses avoided. Attestiv’s approach underlines that **every fake photo caught is money saved**, and they claim even a

modest deployment can pay back quickly (one analysis suggests well-built fraud detection can pay for itself in under 7 months by savings <sup>63</sup> ).

- **Truepic (and similar authenticity tools):** Truepic provides a **secure camera technology** that ensures photos are **verified at capture** (with cryptographic stamps, secure timestamps, and location verification). While not a tampering “detection” tool per se, it’s used in insurance to **prevent tampering** – for example, some insurers require claim photos to be taken through a Truepic-powered app which guarantees the image hasn’t been edited and is time/location-authentic. This approach circumvents the need for after-the-fact forensics by establishing a chain of trust for the images. Truepic’s system can immediately flag if an image has been manipulated or is a screen upload instead of a live photo. Such solutions are more popular in **underwriting inspections and low-touch claims:** e.g., verifying a car’s condition at policy inception or fast-tracking a claim if the images come certified. In Europe, insurers concerned with fraud may adopt this for customer self-service claim apps. The **trade-off** is requiring user compliance (policyholders must use the special app), which can affect user experience. Cost-wise, this is often provided as an SDK or service license to the insurer.
- **Others:** There are numerous other players and in-house integrators. For example, Spain’s **Bdeo** and France’s **Delfos** offer AI damage assessment that can include fraud flags (like detecting if an image is of a model car toy or a screen photo). Companies like **Microsoft (Azure)** and **Google** are also adding AI vision services that could be leveraged to detect image anomalies, although these are not insurance-specific. Additionally, forensic software like **Amped Authenticate** (used in law enforcement) can be used by insurers’ investigation units to do deep forensic analysis on suspicious images, though those are more manual tools than automated solutions.

**Cost & Deployment:** Commercial solutions typically are offered as **cloud-based services or on-premise software**. Insurers can integrate via API calls (sending images to the vendor’s engine for analysis) or install the solution in-house for data privacy. **Pricing models** vary: some charge **per image or per claim screened**, others a flat **annual license** or subscription based on the insurer’s size and volume. For instance, a SaaS might charge a few cents to a few dollars per image analysis depending on volume, whereas an enterprise license for a full platform could run into hundreds of thousands of euros annually for a large insurer. Despite the upfront cost, the **ROI is usually strong** – by preventing even a handful of large fraudulent payouts, the solution can pay for itself. Vendors often cite ROI figures >200% and case studies where millions in fraud were avoided within the first year of implementation. Moreover, **automation savings** (reduced manual reviews, faster claims handling) add to the benefits. In summary, **COTS solutions offer state-of-the-art tech “out of the box”** – ideal for insurers who need a fast, proven way to combat image fraud. They continuously update their models (shared across clients) and handle maintenance, which is valuable given how quickly AI-fraud tactics evolve.

## In-House Solutions & Strategies

Some large insurance organizations pursue an **in-house approach** to image fraud detection, or a hybrid of in-house and vendor tools:

- **Custom AI Models:** Insurers with strong data science teams (or innovation labs) sometimes develop their own **machine learning models** to detect tampered images. For example, Allianz UK built an internal ML tool called **“Incognito”** to flag potentially fraudulent claims (not limited to images, but including them) <sup>64</sup> . This system scans incoming claims and routes suspicious ones to fraud experts, and in its early use it **saved about £1.7 million** by catching fraud that would have been paid out <sup>65</sup> . An in-house model can be trained on the insurer’s historical claims and known fraud cases, potentially giving it an edge on company-specific fraud patterns.

It also allows the insurer to fully control the system and customize it to their workflows. **However, building such models is non-trivial** – it requires collecting a large dataset of both legitimate and fraudulent images, labeling them, and employing skilled ML engineers. Development costs can easily range from **\$200k to \$800k+** depending on complexity <sup>63</sup>, and ongoing maintenance is needed to keep up with new fraud trends. Only the largest insurers tend to have the scale for this investment.

- **Business Rules & Analytics:** Many insurers historically relied on a set of **business rules** (e.g. “flag any claim with missing metadata” or “flag if repair estimate photo is a screenshot”). These can be developed in-house and integrated into claim systems. While rules alone miss sophisticated fakes, insurers often still use them in combination with AI scores. The advantage is that adjusters can understand a rule-based flag (it’s explainable), and they can tweak rules quickly if they see new fraud modus operandi. Modern fraud systems allow insurers to input their own custom rules on top of AI outputs – for example, auto-flagging a claim if an image fails the AI authenticity check *and* the claimant is a new customer with high claim value, etc. This combination of **in-house domain knowledge with vendor AI** often yields the best results <sup>42</sup>.
- **Internal Fraud Teams & Tools:** Even with automated detection, insurers maintain **Special Investigation Units (SIUs)** that handle suspected fraud. These teams may use specialized forensic tools internally. For instance, an investigator might use reverse image search manually, check the claimant’s social media for the same photos, or use software like Amped or Photoshop to examine image layers and metadata in detail. Insurers also participate in **industry data-sharing** (in Germany, the GDV has databases; in the UK, the Insurance Fraud Bureau and databases like MID for motor). These in-house efforts complement automated tools. Essentially, **in-house processes set the strategy and policies** – e.g. deciding that any AI-flagged image triggers a secondary review rather than outright denial – to balance fraud control with customer service <sup>48</sup>.
- **Hybrid Integration:** The prevalent model in Europe is a **hybrid**: use COTS AI services for the heavy analysis, but integrate them tightly with in-house systems and data. Insurers often embed an image-fraud API into their claim management system so that when adjusters or even customers upload photos, the system instantly returns a credibility score or alert. The adjuster sees something like “Image authenticity check: Failed – duplicated regions detected” and can then investigate further. The insurer’s in-house team might also feed feedback loops – e.g. if an adjuster confirms a fraud case that the AI flagged, that data can be used to retrain models (with the vendor or internal model). **Data privacy and sovereignty** influence decisions here: German insurers, for instance, are very sensitive about customer data under GDPR. Some prefer on-premise deployment of AI or at least guarantees that images won’t be stored by a vendor <sup>46</sup>. This can tilt an insurer towards either building in-house or choosing a vendor that allows a private instance. VAARHAFT’s emphasis on no customer data storage and not training on the insurer’s data without consent is a nod to this concern [28timage] .

**When to build in-house vs buy:** Generally, **in-house development is recommended only for specific aspects** or for insurers with unique requirements:

- **Leverage COTS for core detection:** The consensus is that the **specialized AI models (for detecting pixel manipulation, deepfakes, etc.) are best provided by expert vendors** who focus on that technology. They can aggregate learning across many fraud examples and respond quickly to new AI threats (e.g. new deepfake generators) in a way that would be hard for a single insurer’s IT team to match <sup>36</sup>. Buying a proven image forensics solution also speeds up deployment – critical given how fast fraud is evolving.

- **Keep strategic data & decisions in-house:** Insurers should maintain control of their **fraud decision rules, thresholds, and integration with claim workflows**. For example, deciding what fraud score triggers an investigation, or combining image-fraud signals with other fraud indicators (like claim history) – those logic and strategy elements are typically done in-house or in a custom manner. The **knowledge of an insurer's portfolio** (types of claims, customer base, common fraud modus operandi in their market) is valuable in tuning the system. An insurer might develop an in-house dashboard that shows all alerts (image fraud, claim analytics, etc.) in one place for their investigators.
- **Cost considerations:** For many mid-sized insurers, the cost of building and continuously updating a sophisticated image forensics AI from scratch would outweigh the cost of licensing one. The market solutions are spreading those R&D costs over many clients. However, very large insurance groups might pursue partial in-house builds to avoid per-transaction fees and to own the IP long-term. We see some big players like Allianz investing in AI fraud labs. Still, even they often use a blend (Allianz, for instance, might develop a core fraud engine but also use third-party tools for specific tasks like document fraud or image deepfake detection). A published guide suggests fraud detection software development can easily run hundreds of thousands of dollars, but the **ROI of effective fraud detection can exceed 200%** with payback in under a year <sup>63</sup>. That implies that whether bought or built, a working solution basically **pays for itself by cutting fraud losses**.
- **Collaboration and Data Sharing:** One advantage of vendor solutions is **network effects** – e.g. if multiple insurers feed into a system, it can catch cross-insurer scams (as Vaarhaft does) or identify emerging fraud patterns faster. In-house systems would miss those broader patterns unless insurers collaborate. Europe is gradually moving toward more cross-border and cross-company cooperation on fraud <sup>66</sup> <sup>1</sup>, but data privacy and competitive concerns make it tricky. Using a neutral third-party platform can sometimes facilitate this sharing (since the vendor can act as a custodian of pooled fraud data). Insurers must weigh this benefit against handing data to a third party. In practice, many resolve it by contractual agreements that the data is only used for their service and anonymized if aggregated.

## Use Case & Trade-off Summary

- **Accuracy vs. Explainability:** AI deepfake detectors can be very accurate but are often “black boxes.” In-house staff and regulators may demand explanations for why a claim was flagged (especially if denying a claim). COTS vendors are adding features to highlight the manipulated area on the image or provide reason codes (e.g. “Photo appears digitally altered in region X”), which helps bridge this gap <sup>39</sup>. Insurers may also set thresholds such that borderline cases are manually reviewed to avoid false denials.
- **Automation vs. Human Oversight:** The ultimate goal is **straight-through processing for legitimate claims and automated blocking of fraudulent ones**. However, a false positive has a high cost (customer dissatisfaction, potential legal issues), so most insurers still keep a human-in-the-loop for fraud decisions. The trade-off is tuning the system to catch maximum fraud without overburdening adjusters with too many false alerts. Multi-layered approaches (metadata + AI + context) help reduce noise by requiring multiple red flags before interruption <sup>67</sup>. Over time, as confidence in AI grows, we might see more fully automated denial of clearly fake claims, but currently in Europe the norm is to use these tools to **assist investigators**, not replace them entirely <sup>68</sup> <sup>45</sup>.

- **Cost of fraud vs. cost of solution:** In Germany, since ~€3 billion of motor claims a year might be fraudulent <sup>4</sup>, even a 10% improvement in detection could save €300 million industry-wide – a huge win. The cost of implementing an AI solution is relatively small compared to this. **Solution costs** can be viewed per claim: for example, if manual review of a suspicious claim costs €50–€100 in investigator time, an automated check that costs a few euros (or less) per claim is very attractive <sup>69</sup> <sup>70</sup>. Many vendors thus price their services to be a fraction of the average fraud savings per claim. Insurers will compare the **in-house cost** of doing the same (both in IT expenses and in missed fraud) when justifying purchases. Given the high fraud costs, most see these tools as **high ROI investments** rather than expenses.

## Conclusion: Post-Incident Image Fraud Detection with AI

In the European insurance industry – and particularly in Germany – **image tampering detection has become a critical component of claims processing**. The surge of digital fraud (from edited accident photos to AI-generated damage scenes) demands state-of-the-art defenses. **ML/AI plays an indispensable role**: it enables real-time, scalable scrutiny of every photo, catching subtle fakes that humans would overlook <sup>71</sup> <sup>72</sup>. Equally important is integrating these tools into the claims workflow (at FNOL or early in assessment) so that **fraudulent claims are flagged before payment** without grinding the process to a halt <sup>45</sup>. Europe's insurers are embracing a mix of **COTS solutions and in-house strategies** to achieve this. Commercial solutions bring specialized technology – continuously updated to counter new AI threats – while in-house teams bring domain knowledge and oversight to deploy these tools effectively and in compliance with local regulations (GDPR, etc.).

Going forward, we can expect:

- **Broader adoption in Germany/Europe** as even mid-sized insurers adopt cloud-based fraud detection APIs (the market growth figures show high adoption rates year over year). Fraud rings often operate across borders, so a more unified European approach (via platforms or insurer alliances) may emerge <sup>66</sup>.
- **More AI advancements** like ensemble models that combine image analysis with voice analytics or other fraud signals. (E.g., some insurers also use voice-stress AI for phone claims – a complementary tech, though outside our image scope.) The arms race with fraudsters will continue, with **generative AI both a threat and a tool** for the industry <sup>37</sup> <sup>38</sup>.
- **Cost efficiency improvements:** as these solutions mature, the cost per screening may drop, and usage will become ubiquitous (just as spam email filters are standard today). The focus will shift from just detection to **deterrence** – if fraudsters know that insurers can detect photoshopped or AI images reliably, it may discourage attempts in the first place.

In summary, **post-incident image fraud detection in car claims is now a high-tech endeavor**. Insurers in Germany and across Europe recognize that trusting digital images blindly is not viable in 2025. By leveraging sophisticated AI-driven tools (often via COTS products) and blending them with in-house fraud management, they are improving their fraud catch rates, protecting their bottom line, and keeping the claims process fast for honest customers <sup>24</sup> <sup>73</sup>. The investment in these technologies is driven by clear ROI – not only in euros saved from false claims, but in preserving the overall **trust and credibility** of an increasingly digital claims landscape.

**References:** (European insurance industry sources and tech provider documentation have been used to compile this analysis, with particular focus on Germany.)

- <sup>3</sup> <sup>4</sup> GDV (German Insurers Association) via XPRIMM – *Fraud losses in Germany (~10% of claims, ~€6B annually, half in motor)*.

- 1 Insurance Europe – *Estimated 10% of claims costs lost to fraud across Europe.*
  - 74 9 The Guardian – *Rising cases of fake damage photos in UK (300% increase in edited images; "shallowfakes" an emerging threat).*
  - 36 39 InsureTech Connect Europe (Vaarhaft) – *Vaarhaft's AI approach (deep learning, fast adaptation) and features like cross-insurer duplicate image check.*
  - 44 47 Inaza (Insurtech) – *Layers of AI image analysis (pixel, metadata, semantic) and current limitations (false positives, evolving tricks).*
  - 72 32 Shift Technology – *Use of ELA for splicing detection and shift to Vision Transformers for AI-generated image detection (state-of-art ML).*
  - 45 51 FRISS – *Multi-layered fraud defense: forensic tools, workflow integration, and Media Check solution capabilities (metadata, anomaly, reuse analysis integrated into claims).*
  - 53 54 FRISS Press Release – *FRISS & Verisk collaboration bringing AI-driven Digital Media Forensics to EU insurers (analyze each photo for anomalies, trace file origin, detect edits).*
  - 64 65 Allianz UK – *In-house "Incognito" ML tool flags fraudulent claims, saving £1.7M; AI used to stay ahead of fraud rings.*
  - 63 ScienceSoft – *Cost of developing fraud detection software (\$200k-\$800k+) and rapid ROI (<7 months) due to savings.*
  - 60 61 Attestiv – *Emphasis on AI catching fake media to reduce manual reviews and quantify fraud losses (ROI calculator for insurance media fraud).*
  - 18 75 FRISS Blog – *Examples of image-based fraud: duplicated damages ("shallowfakes"), reused stock images, AI-generated fake scenes.*
  - 20 The Guardian – *Example of fraud: using salvage car photos with fake plates to claim nonexistent crashes ("people can now create a fraudulent claim entirely from behind their computer").*
  - 49 FRISS – *Contextual analysis reduces false positives (e.g. compare metadata timing with claim info, route high-risk anomalies to human SIU for review).*
  - 46 Vaarhaft – *Real-time API checks in seconds; state-of-art tech, GDPR compliant, customizable per client.*
  - 38 Shift – *"Fighting fire with fire": using GenAI to detect AI-fraud, since fake evidence is easy to produce for average person.*
- 

1 66 Annual Report 2023-2024: Fraud

<https://www.insuranceeurope.eu/downloads/ar-2024-fraud/Fraud.pdf>

2 8 13 24 36 39 46 56 57 58 VAARHAFT: AI detection of fake claim images to fight insurance fraud - ITC Europe 2025

<https://europe.insuretechconnect.com/news-articles/vaarhaft-ai-detection-fake-claim-images-fight-insurance-fraud>

3 4 GDV: Costs of fraudulent claims is increasing, in line with the increasing benefits in P&C insurance - Insurance -

<https://www.xprimm.com/GDV-Costs-of-fraudulent-claims-is-increasing-in-line-with-the-increasing-benefits-in-P-C-insurance-articol-124-21762.htm>

5 6 7 9 16 19 20 21 74 Fraudsters editing vehicle photos to add fake damage in UK insurance scam | Insurance industry | The Guardian

<https://www.theguardian.com/business/article/2024/may/02/car-insurance-scam-fake-damaged-added-photos-manipulated>

10 11 Insurance Fraud Detection Market Report 2025-2034 | Trends

<https://www.thebusinessresearchcompany.com/report/insurance-fraud-detection-global-market-report>

12 From "Traditional AI" to "Generative AI": Implications for the ... - EIOPA

[https://www.eiopa.europa.eu/publications/traditional-ai-generative-ai-implications-insurance-sector\\_en](https://www.eiopa.europa.eu/publications/traditional-ai-generative-ai-implications-insurance-sector_en)

14 15 Insurance Fraud Detection Market in Germany

<https://www.lucintel.com/insurance-fraud-detection-market-in-germany.aspx>

17 23 25 33 34 35 43 44 47 50 Detecting Tampered Images in Insurance Claims | Inaza

<https://www.inaza.com/blog/detecting-tampered-images-in-insurance-claims>

18 29 30 42 45 48 49 51 55 67 68 73 75 Digital Deceit: The Rising Threat of Image Alteration

Fraud - FRISS

<https://www.friiss.com/blog/the-rising-threat-of-image-alteration-fraud>

22 26 27 28 31 32 37 38 71 72 Shift Insurance Perspectives: The Fraudulent Image Analysis Edition

<https://www.shift-technology.com/resources/reports-and-insights/shift-insurance-perspectives-the-fraudulent-image-analysis-edition>

40 AI in Action: Uncovering reused photo fraud with Shift's AI-based ...

<https://www.shift-technology.com/resources/case-studies/uncovering-reused-photo-fraud>

41 Shift Technology helps hundreds of insurers fight claims fraud using ...

<https://www.microsoft.com/en/customers/story/23202-shift-technology-azure-ai-vision>

52 53 54 FRISS and Verisk Claims UK Join Forces to Combat Digital Media Fraud - FRISS

<https://www.friiss.com/press/friiss-and-verisk-claims-uk-join-forces-to-combat-digital-media-fraud>

59 Infographic: Using advanced AI and insurance expertise to find and ...

<https://www.shift-technology.com/resources/reports-and-insights/ai-and-insurance-expertise-for-detecting-document-fraud-in-fwa-cases>

60 61 62 Introducing Attestiv's ROI Calculator: See the Real Cost of Insurance Fraud — And How to Beat It - Attestiv

<https://attestiv.com/introducing-attestivs-roi-calculator-see-the-real-cost-of-insurance-fraud-and-how-to-beat-it/>

63 Insurance Fraud Detection: A Guide to 200%+ ROI - ScienceSoft

<https://www.scnsoft.com/insurance/fraud-detection>

64 Combating Insurance Frauds Through Data Management and AI

<https://www.infoverity.com/en/blog/combating-insurance-frauds-through-advanced-data-management-and-ai/>

65 AI Meets Insurance: Revolutionizing Claims with Automation - Aranca

<https://www.aranca.com/knowledge-library/articles/business-research/ai-meets-insurance-revolutionizing-claims-with-automation>

69 Insurance claims | Resistant AI

<https://resistant.ai/use-case/insurance-fraud>

70 Anti-Fraud Claims Solutions | Verisk

<https://www.verisk.com/solutions/claims/anti-fraud/>